

5. Regularization

ALLab
Hanyang Univ.


오늘 실습 내용

1. Overfitting
2. L1 L2 Regularization
3. Dropout
4. Normalization

1. Overfitting

• Overfitting이란?

- 한 데이터셋에만 지나치게 최적화된 상태
- 아래 그래프처럼 학습 데이터에 대해서는 오차가 감소하지만 실제 데이터에 대해서는 오차가 증가하는 지점이 존재할 수 있음
- 즉, overfitting은 학습데이터에 대해 과하게 학습하여 실제 데이터에 대한 오차가 증가할 경우 발생



Overfitting

- **Overfitting** 을 완화시키는 방법은?
- Training Data 를 늘린다.
- Regularization
- Dropout
- Normalization

2. L1 L2 Regularization

Regularization

- 모델의 파라미터 확인하는 방법
 - Use `named_parameters()` or `parameters()`
 - https://pytorch.org/docs/stable/generated/torch.nn.Module.html?highlight=named_parameters#torch.nn.Module.named_parameters

```
class LogisticRegression(nn.Module):
    def __init__(self, x_in, x_out):
        super(LogisticRegression, self).__init__()
        self.linear = nn.Linear(x_in, x_out)
        self.activation = nn.Sigmoid()
    def forward(self, x):
        z = self.linear(x)
        a = self.activation(z)
        return a
```



```
for name, param in model.named_parameters():
    print('=====')
    print(name)
    print(param.shape)
    print(param)

=====
linear.weight
torch.Size([1, 2])
Parameter containing:
tensor([[ 0.0162, -0.1808]], requires_grad=True)
=====
linear.bias
torch.Size([1])
Parameter containing:
tensor([-0.0776], requires_grad=True)
```

Regularization

- L1 loss in LogisticRegression

`reg = model.linear.weight.abs().sum()`

$$\|w\|_1 = \sum_{j=1}^n |w_j|$$

```
print(model.linear.weight)
print(model.linear.weight.abs().sum())
```

```
Parameter containing:
tensor([[ 0.0162, -0.1808]], requires_grad=True)
tensor(0.1970, grad_fn=<SumBackward0>)
```


Regularization

- L2 loss in LogisticRegression

`reg = model.linear.weight.pow(2.0).sum()`

$$\|w\|_2^2 = \sum_{j=1}^n w_j^2 = w^T w$$

```
print(model.linear.weight)
print(model.linear.weight.pow(2.0).sum())
```

```
Parameter containing:
tensor([[ 0.0162, -0.1808]], requires_grad=True)
tensor(0.0330, grad_fn=<SumBackward0>)
```

Regularization

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)}) + \frac{\lambda}{2m} \|w\|_2^2$$

- L1 loss

```
reg = model.linear.weight.abs().sum()
```

```
loss = loss + lambda * reg/total_num/2.
```

- L2 loss

```
L2_norm = model.linear.weight.pow(2.0).sum()
```

```
loss = loss + lambda * reg/total_num/2.
```

~~Logistic Regression의 overfitting 확인하거나 성능 확인하는 test set~~

Cifar-10 Data

Cifar-10

- **Cifar-10 Dataset**
 - 10개의 클래스로 분류된 RGB image
 - 6000 32x32 per class. 50000 training images, 10000 test images.

airplane



automobile



bird



cat




deer



dog



frog



horse



ship




truck




Cifar-10

- RGB image
 - View 이용해서 shape 조정



3D data → 1D data




Cifar-10

- Image Data in Simple Neural Network
 - 실제로 view가 잘 실행되는지 확인

img

shape : (3,3,3)

**new_img**

shape : (27)



```
img = torch.FloatTensor([ [.392,.482,.576],[.478,.639,.241], [.580,.790,.543]],
                          [.169,.263,.376],[.263,.442,.823],[.373,.602,.165]],
                          [.306,.376,.451],[.376,.478,.561],[.443,.569,.674]] )
```

```
print(img.shape)
new_img = img.view(3*3*3)
print(new_img.shape)
print(new_img)
print(new_img[3*3*3-3:3*3*3])
print(img[2,2,:])
```

```
torch.Size([3, 3, 3])
torch.Size([27])
tensor([0.3920, 0.4820, 0.5760, 0.4780, 0.6390, 0.2410, 0.5800, 0.7900, 0.5430,
        0.1690, 0.2630, 0.3760, 0.2630, 0.4420, 0.8230, 0.3730, 0.6020, 0.1650,
        0.3060, 0.3760, 0.4510, 0.3760, 0.4780, 0.5610, 0.4430, 0.5690, 0.6740])
tensor([0.4430, 0.5690, 0.6740])
tensor([0.4430, 0.5690, 0.6740])
```


Regularization

- Random Seed 고정
 - 모델 weight가 생성될 때마다 random하게 생성됨
 - 성능 비교를 하기위해 값을 고정하는 것이 좋다.
 - 어떤 random한 값에서는 좋게 나오고 다른 값에서는 나쁘게 나올 수 있음

```
import torch
import torch.nn as nn
import torch.optim as optim
```

```
torch.manual_seed(0)
torch.cuda.manual_seed(0)
torch.cuda.manual_seed_all(0)
```

```
print(model.linear1.weight[0,1])
```

```
tensor(-0.0298, device='cuda:0', grad_fn=<SelectBackward0>)
```


Regularization

- Cifar-10 Dataset
 - Model 구조
 - Input size : $32 \times 32 \times 3$
 - Output size : 10

```
class Model(nn.Module):  
    def __init__(self):  
        super(Model, self).__init__()  
        self.linear1 = nn.Linear(32*32*3, 256)  
        self.linear2 = nn.Linear(256, 128)  
        self.linear3 = nn.Linear(128, 10)  
  
        self.activation = nn.Sigmoid()  
  
    def forward(self, x):  
        z1 = self.linear1(x)  
        a1 = self.activation(z1)  
  
        z2 = self.linear2(a1)  
        a2 = self.activation(z2)  
  
        z3 = self.linear3(a2)  
  
        return z3
```

Regularization

- 실제로 학습이 overfitting인지 확인
 - Matplotlib 활용
 - detach()
 - 기존 텐서 복사
 - cpu()
 - GPU에 있는 tensor CPU로 이동

```
epochs = 70
train_avg_costs = []
test_avg_costs = []

test_total_batch = len(test_dataloader)
total_batch_num = len(train_dataloader)

for epoch in range(epochs):
    avg_cost = 0
    model.train()
    for b_x, b_y in train_dataloader:
        b_x = b_x.view(-1, 32*32*3).to(device)
        logits = model(b_x) # forward propagation
        loss = criterion(logits, b_y.to(device)) # get cost

        optimizer.zero_grad()
        loss.backward() # backward propagation
        optimizer.step() # update parameters

    avg_cost += loss / total_batch_num
    train_avg_costs.append(avg_cost.detach().cpu())
    print('Epoch : {} / {}, cost : {}'.format(epoch+1, epochs, avg_cost))


    test_avg_cost=0
    model.eval()
    for b_x, b_y in test_dataloader:
        b_x = b_x.view(-1, 32*32*3).to(device)
        with torch.no_grad():
            logits = model(b_x)
            test_loss = criterion(logits, b_y.to(device)) # get cost
        test_avg_cost += test_loss / test_total_batch

    test_avg_costs.append(test_avg_cost.detach().cpu())
```

Regularization

- 실제로 학습이 overfitting인지 확인
 - Matplotlib 활용
 - Test Loss 증가 확인

```
import matplotlib.pyplot as plt
import numpy as np
epoch = range(epochs)
plt.plot(epoch, train_avg_costs, 'r-')
plt.plot(epoch, test_avg_costs, 'b-')
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend(['train', 'test'])
plt.show()
```



Regularization

- 전체 코드 1/3

```
[ ] torch.manual_seed(0)
    torch.cuda.manual_seed(0)
    torch.cuda.manual_seed_all(0)
```

```
[ ] if torch.cuda.is_available():
    device = torch.device('cuda')
else:
    device = torch.device('cpu')
```

```
▶ import torchvision
import torchvision.transforms as transforms

train_dataset = torchvision.datasets.CIFAR10(root="CIFAR10/",
                                             train=True,
                                             transform=transforms.ToTensor(),
                                             download=True)
test_dataset = torchvision.datasets.CIFAR10(root="CIFAR10/",
                                             train=False,
                                             transform=transforms.ToTensor(),
                                             download=True)
```

```
▶ batch_size = 128
```

```
train_dataloader = torch.utils.data.DataLoader(train_dataset, batch_size=batch_size)
test_dataloader = torch.utils.data.DataLoader(test_dataset, batch_size=batch_size)
```

```
[ ] class Model(nn.Module):
    def __init__(self, drop_prob):
        super(Model, self).__init__()
        self.linear1 = nn.Linear(32*32*3, 256)
        self.linear2 = nn.Linear(256, 128)
        self.linear3 = nn.Linear(128, 10)

        self.activation = nn.Sigmoid()

    def forward(self, x):
        z1 = self.linear1(x)
        a1 = self.activation(z1)

        z2 = self.linear2(a1)
        a2 = self.activation(z2)

        z3 = self.linear3(a2)

        return z3
```

Regularization

- 전체 코드 2/3

```
[ ] model = Model().to(device).train()

[ ] optimizer = optim.SGD(model.parameters(), lr=1) # set optimizer

[ ] criterion = nn.CrossEntropyLoss()

[ ] epochs = 70

train_avg_costs = []
test_avg_costs = []

test_total_batch = len(test_dataloader)
total_batch_num = len(train_dataloader)

for epoch in range(epochs):
    avg_cost = 0
    model.train()
    for b_x, b_y in train_dataloader:
        b_x = b_x.view(-1, 32*32*3).to(device)
        logits = model(b_x) # forward propagation
        loss = criterion(logits, b_y.to(device)) # get cost

        optimizer.zero_grad()
        loss.backward() # backward propagation
        optimizer.step() # update parameters

    avg_cost += loss / total_batch_num
    train_avg_costs.append(avg_cost.detach())
    print('Epoch : {} / {}, cost : {}'.format(epoch+1, epochs, avg_cost))

    test_avg_cost=0
    model.eval()
    for b_x, b_y in test_dataloader:
        b_x = b_x.view(-1, 32*32*3).to(device)
        with torch.no_grad():
            logits = model(b_x)
            test_loss = criterion(logits, b_y.to(device)) # get cost
            test_avg_cost += test_loss / test_total_batch

    test_avg_costs.append(test_avg_cost.detach())
```

Regularization

- 전체 코드 3/3

```
import matplotlib.pyplot as plt
import numpy as np
epoch = range(epochs)
plt.plot(epoch, train_avg_costs, 'r-')
plt.plot(epoch, test_avg_costs, 'b-')
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend(['train', 'test'])
plt.show()
```

Regularization

- L2 Regularization

Regularization for Neural Network

$$J(W^{[1]}, b^{[1]}, \dots, W^{[L]}, b^{[L]}) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)}) + \frac{\lambda}{2m} \sum_{l=1}^L \|W^{[l]}\|_F^2$$

Frobenius norm: $\|W^{[l]}\|_F^2 = \sum_{i=1}^{n^{[l]}} \sum_{j=1}^{n^{[l-1]}} (w_{ij}^{[l]})^2$ $W^{[l]}: (n^{[l]}, n^{[l-1]})$

```

▶ epochs = 70
lmbd = 0.003

train_avg_costs = []
test_avg_costs = []

test_total_batch = len(test_dataloader)
total_batch_num = len(train_dataloader)

for epoch in range(epochs):
    avg_cost = 0
    model.train()
    for b_x, b_y in train_dataloader:
        b_x = b_x.view(-1, 32*32*3).to(device)
        logits = model(b_x) # forward propagation
        loss = criterion(logits, b_y.to(device)) # get cost

        reg = model.linear1.weight.pow(2.0).sum()
        reg += model.linear2.weight.pow(2.0).sum()
        reg += model.linear3.weight.pow(2.0).sum()

        loss += lmbd*reg/len(b_x)/2.

    optimizer.zero_grad()
    loss.backward() # backward propagation
    optimizer.step() # update parameters

    avg_cost += loss / total_batch_num
    train_avg_costs.append(avg_cost.detach())
    print('Epoch : {} / {}, cost : {}'.format(epoch+1, epochs, avg_cost))

test_avg_cost=0
model.eval()
for b_x, b_y in test_dataloader:
    b_x = b_x.view(-1, 32*32*3).to(device)
    with torch.no_grad():
        logits = model(b_x)
        test_loss = criterion(logits, b_y.to(device)) # get cost
    test_avg_cost += test_loss / test_total_batch


test_avg_costs.append(test_avg_cost.detach())

```


Regularization

- Loss Comparison

Base Model



L2 Regularization



3. Dropout

Dropout

- Dropout
 - <https://pytorch.org/docs/stable/generated/torch.nn.Dropout.html>

Parameters

- **p** – probability of an element to be zeroed. Default: 0.5
- **inplace** – If set to `True`, will do this operation in-place. Default: `False`

얼만큼 0으로 만들 것의냐의 확률

- `model.eval()`로 모드를 바꿔야 dropout이 작동하지 않음

Dropout

- Dropout

```
class Model(nn.Module):  
    def __init__(self, drop_prob):  
        super(Model, self).__init__()  
        self.linear1 = nn.Linear(32*32*3, 256)  
        self.linear2 = nn.Linear(256, 128)  
        self.linear3 = nn.Linear(128, 10)  
  
        self.dropout = nn.Dropout(drop_prob)  
        self.activation = nn.Sigmoid()  
  
    def forward(self, x):  
        z1 = self.linear1(x)  
        a1 = self.activation(z1)  
        a1 = self.dropout(a1)  
  
        z2 = self.linear2(a1)  
        a2 = self.activation(z2)  
        a2 = self.dropout(a2)  
  
        z3 = self.linear3(a2)  
  
        return z3  
  
model = Model(0.1).to(device).train()
```

4. Normalization

Normalization

- torchvision.transforms.Normalize
 - <https://pytorch.org/vision/stable/generated/torchvision.transforms.Normalize.html#torchvision.transforms.Normalize>

NORMALIZE

`CLASS torchvision.transforms.Normalize(mean, std, inplace=False)` [\[SOURCE\]](#)



Normalize a tensor image with mean and standard deviation. This transform does not support PIL Image. Given mean: `(mean[1], ..., mean[n])` and std: `(std[1], ..., std[n])` for `n` channels, this transform will normalize each channel of the input `torch.*Tensor` i.e., `output[channel] = (input[channel] - mean[channel]) / std[channel]`

1. subtract mean:

$$\mu = \frac{1}{m} \sum_{i=1}^m x^{(i)}$$

$$x^{(i)} := x^{(i)} - \mu$$



2. normalize variance:

$$\sigma_j^2 = \frac{1}{m} \sum_{i=1}^m \{x_j^{(i)}\}^2$$

$$x_j^{(i)} := x_j^{(i)} / \sigma_j$$

오늘 실습 내용

CIFAR-10에 L2 Regularization, Dropout, Normalization 모두 적용해서
test accuracy 확인

→ 만약 예상과 다른 결과가 나온다면 hyperparameter 조정해보기
epoch, learning rate, ...